Clustering by proximity to prototypes

k-means clustering

What is the goal of clustering?

Division of a data set X into k disjoint subsets C_1 … C_k such that objects within each subset are similar and objects in different subsets are dissimilar

Example
of data to be clustered

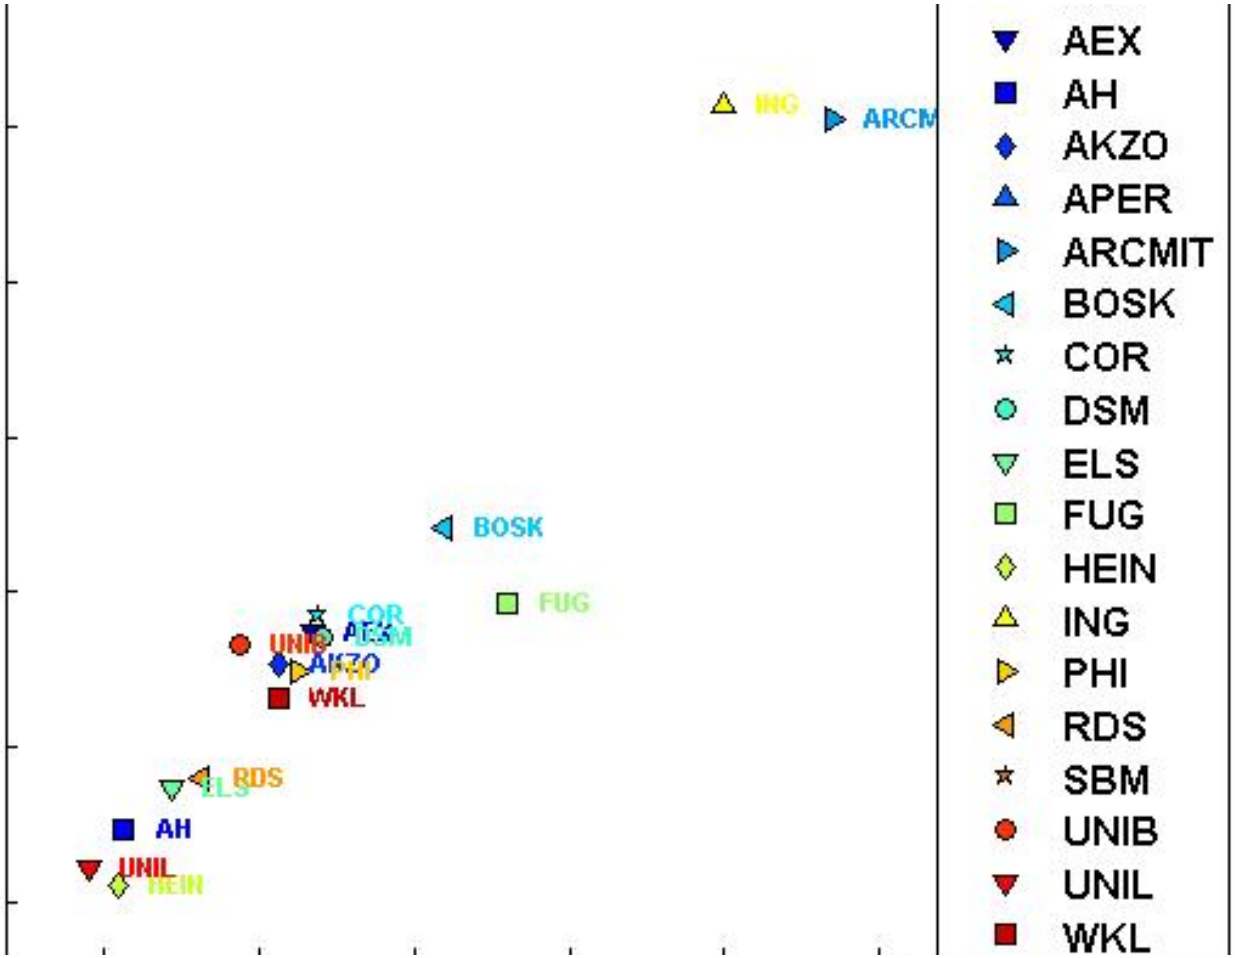

k-means clustering

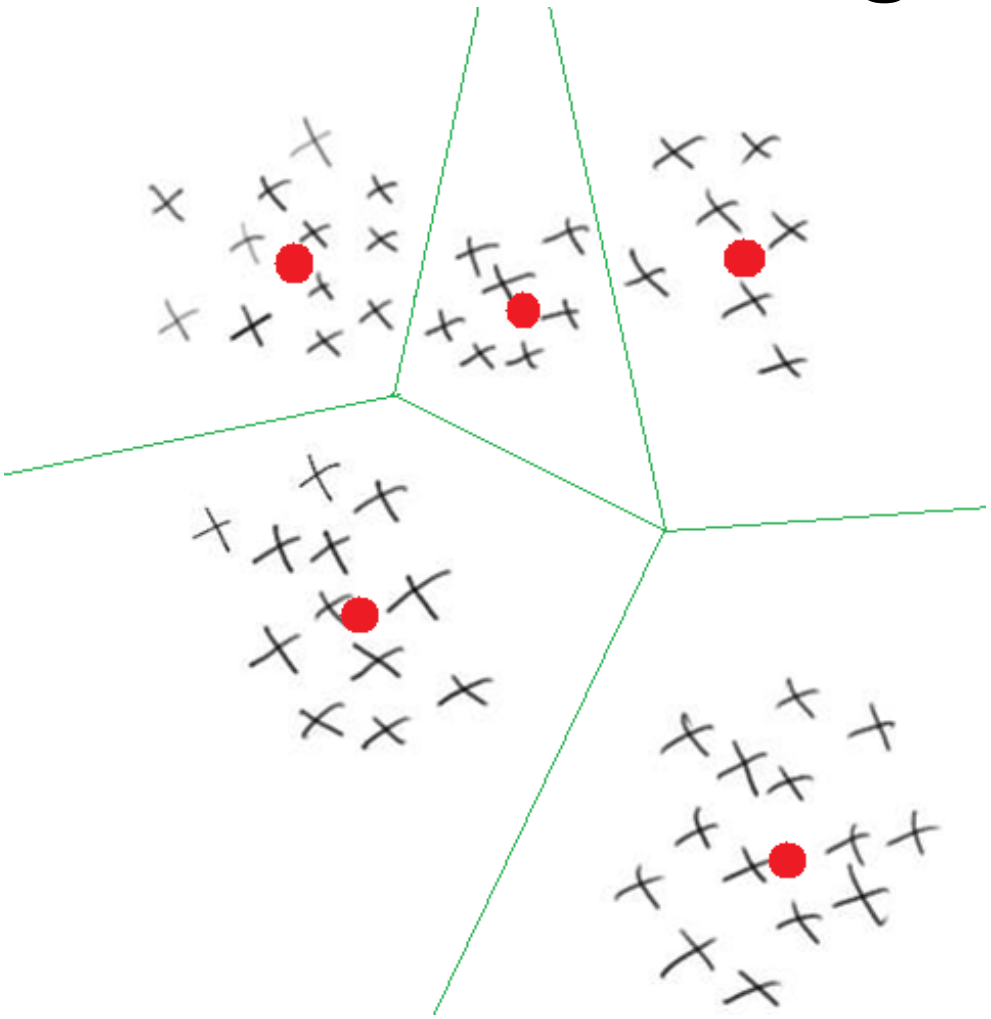

Euclidean-distance, prototype-based clustering:
assign a data point to the nearest prototype

k-means clustering

- given: elements x^j in R^n, number of clusters k
- Goal: find k prototypes μ^i
  that minimize the quantization error

$$J_e = \frac{1}{2}\sum_{\vec{\mu}^i}\sum_{\vec{x}^j \in C(\vec{\mu}^i)} \|\vec{x}^j - \vec{\mu}^i\|^2$$

$C(\mu^i)$ – cluster (subset of X) associated with μ^i
(also called receptive field of μ^i)

Lloyd's algorithm for k-means clustering

1. b**egin initialize** $\mu^1, \mu^2, \ldots, \mu^k$ **(e.g.** take randomly k samples from the data set)
2. **do** *assign data points to nearest* μ^i *(compute* C^i*)*
3. *re-compute* μ^i *as the mean of points in* C^i
4. **until** *no change in* $\mu^1, \mu^2, \ldots, \mu^k$
5. **return** $C^1, C^2, \ldots, C^k$ and $\mu^1, \mu^2, \ldots, \mu^k$
6. **end**

Lloyd's algorithm for k-means clustering

1. b**egin initialize** $\mu^1, \mu^2, \ldots, \mu^k$ **(e.g.** take randomly k samples from the data set)
2. **do** *assign data points to nearest* μ^i *(compute* C^i*)*
3. *re-compute* μ^i *as the mean of points in* C^i
4. **until** *no change in* $\mu^1, \mu^2, \ldots, \mu^k$
5. **return** $C^1, C^2, \ldots, C^k$ and $\mu^1, \mu^2, \ldots, \mu^k$
6. **End**

COMMENT ON HOW TO IMPLEMENT STEP 2 ASSIGNMENT TO CLUSTERS: USE AN INTEGER ATRRAY C such that C(i)=number of cluster to which point i is assigned

Does Lloyd's algorithm converge?

- Yes, in a finite number of steps, because a non-negative cost function (the quantization error) decreases (or remains constant) with each step:

$$J_e = \frac{1}{2}\sum_{\vec{\mu}^i}\sum_{\vec{x}^j \epsilon C(\vec{\mu}^i)} ||\vec{x}^j - \vec{\mu}^i||^2$$

$$\vec{\mu}^i = \frac{1}{n}\sum_{\vec{x}^j \epsilon C(\vec{\mu}^i)} \vec{x}$$

- However, there is no guarantee that a global minimum is reached

Does Lloyd's algorithm converge?

- THE QUANTIZATION ERROR AS A FUNCTION OF THE INTERATION NUMBER MUST DECREASE MONOTONOUSLY
- IF THE QUANTIZATION ERROR SHOWS OSCILLATIONS (goes up and down) THERE MUST BE A BUG IN THE CODE

Intitialisation of k means

- MULTIPLE INITIALIZATIONS, e.g. take data points randomly
- RUN THE K-NN ALGORITHM FOR DIFFERENT INITIALISATIONS AND TAKE THE RESULT FOR WHICH THE QUANITIZATION ERROR IS MINIMUM

HOW TO CHOOSE K?

'ELBOW' METHOD:

1. Run the k-nn algorithm for multiple values of k and for each value of k record the value of the quantization error upon convergence
2. Plot the reached quantization error as a function of k
3. If the plot shows an 'elbow' for a certain k, take that k

Example of k-means clustering

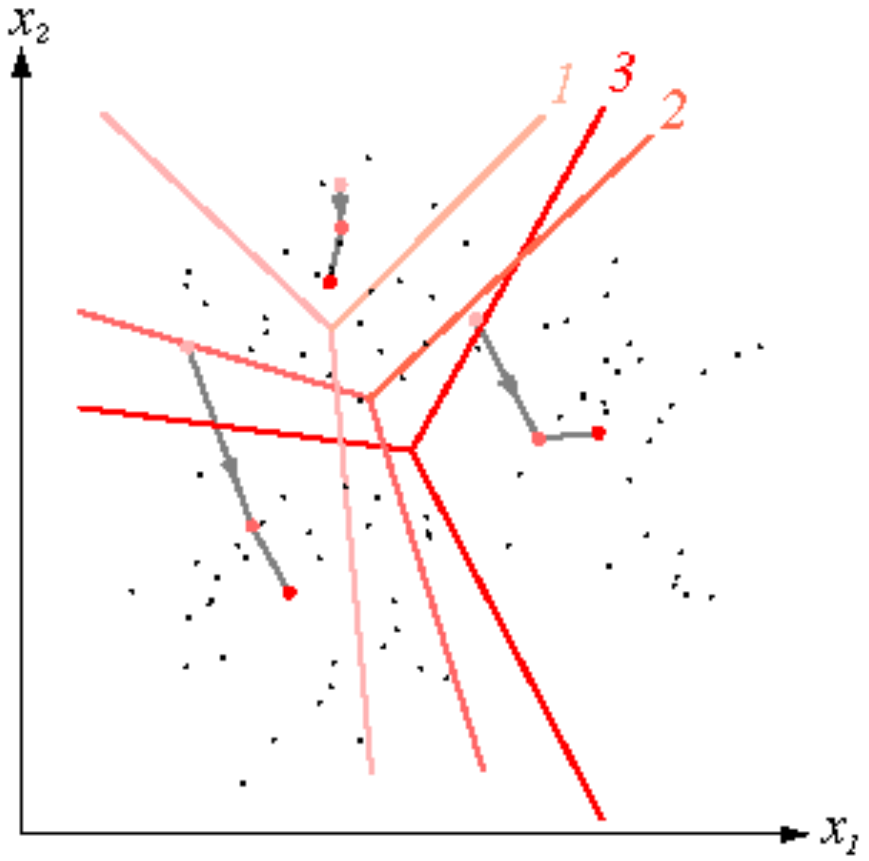

Evolution of the (3) computed means (and Voronoi cells) during 3-means clustering

from Duda, Hart, Stork (2001) Pattern classification

k-means applet

- http://www.cs.rug.nl/~petkov/teaching/PatternRecognition/supplements/k-means/

Iris data

- see UCI repository: http://archive.ics.uci.edu/ml/
  - 150 points
  - 4 dimensional
  - 3 classes

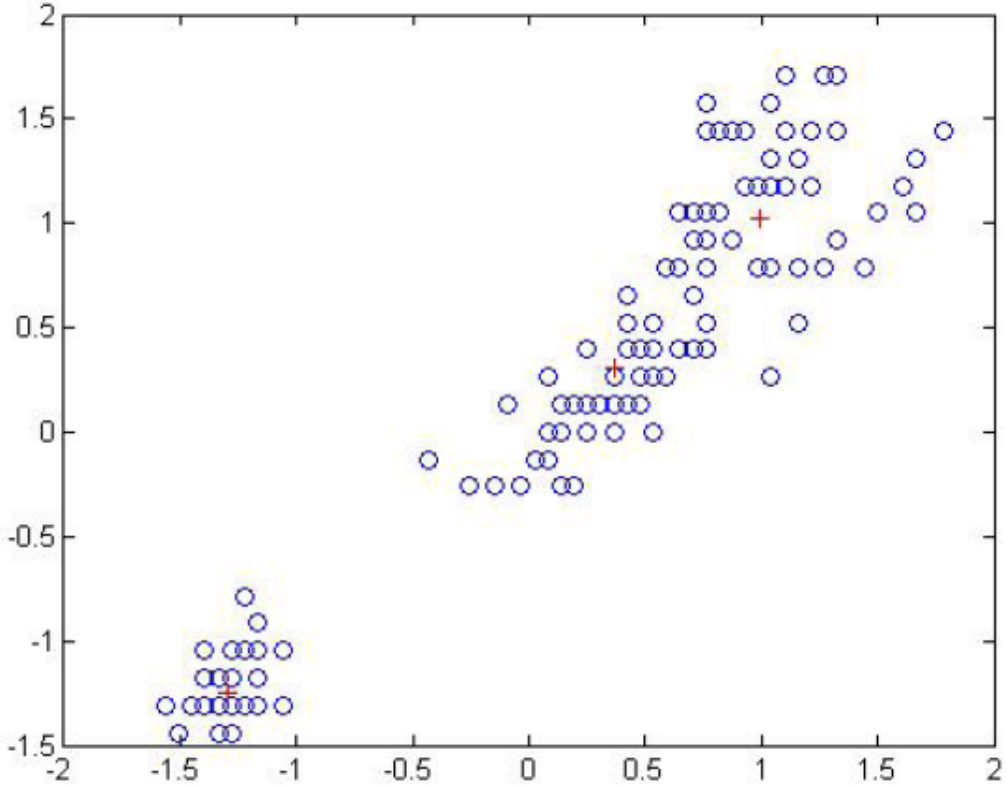

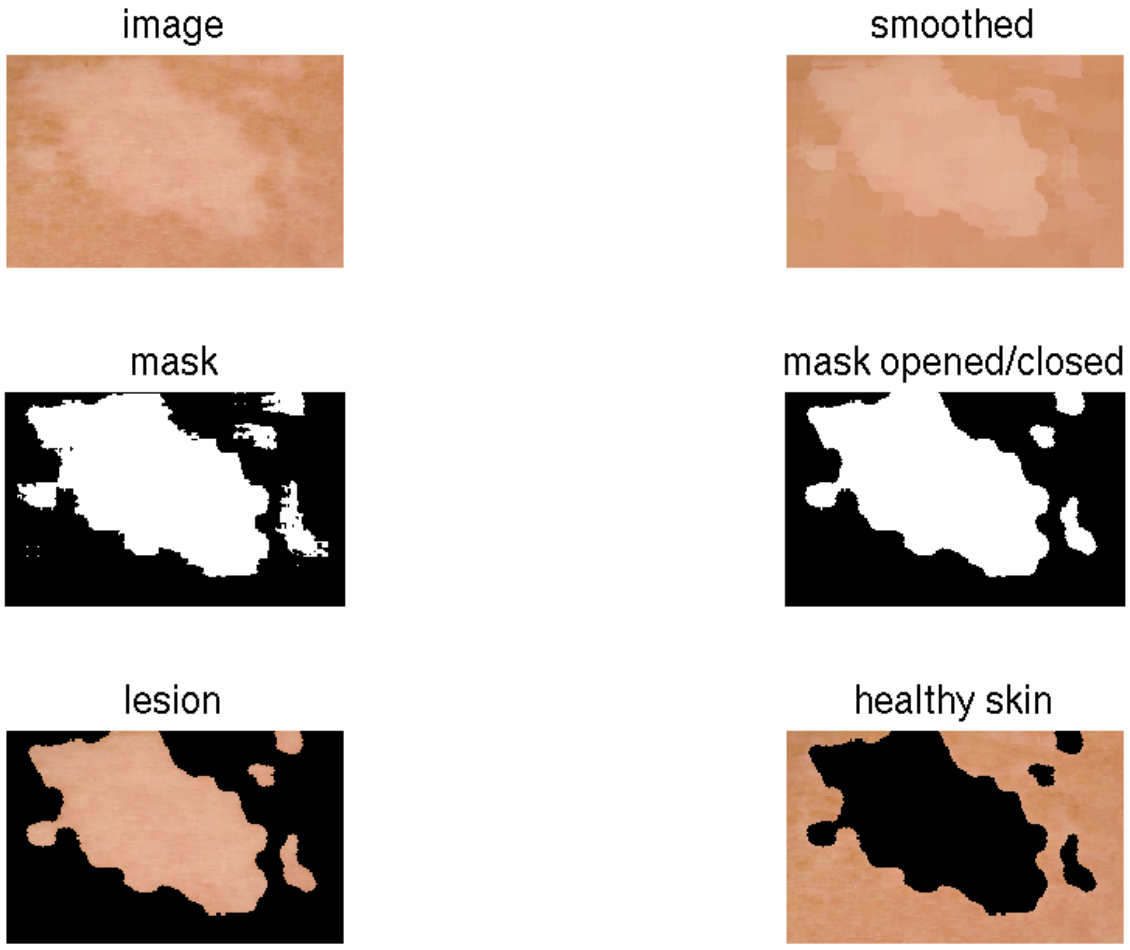

Example of 2-means clustering: a skin image is segmented in two regions of lesion and healthy skin by grouping pixels in two clusters according to their color (result shown in image mask)

Problems with k-means clustering: dead units (poor initialization)

if some prototypes are initialized far away from the input data set, no data points are assigned to them and they are never updated

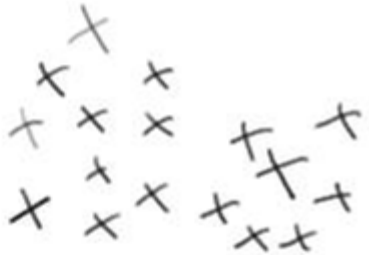

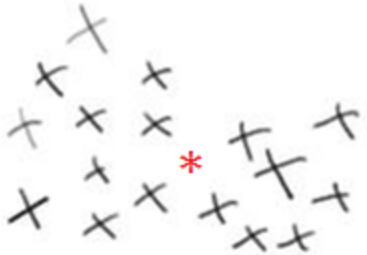

Problems with k-means clustering: non-spherical clusters

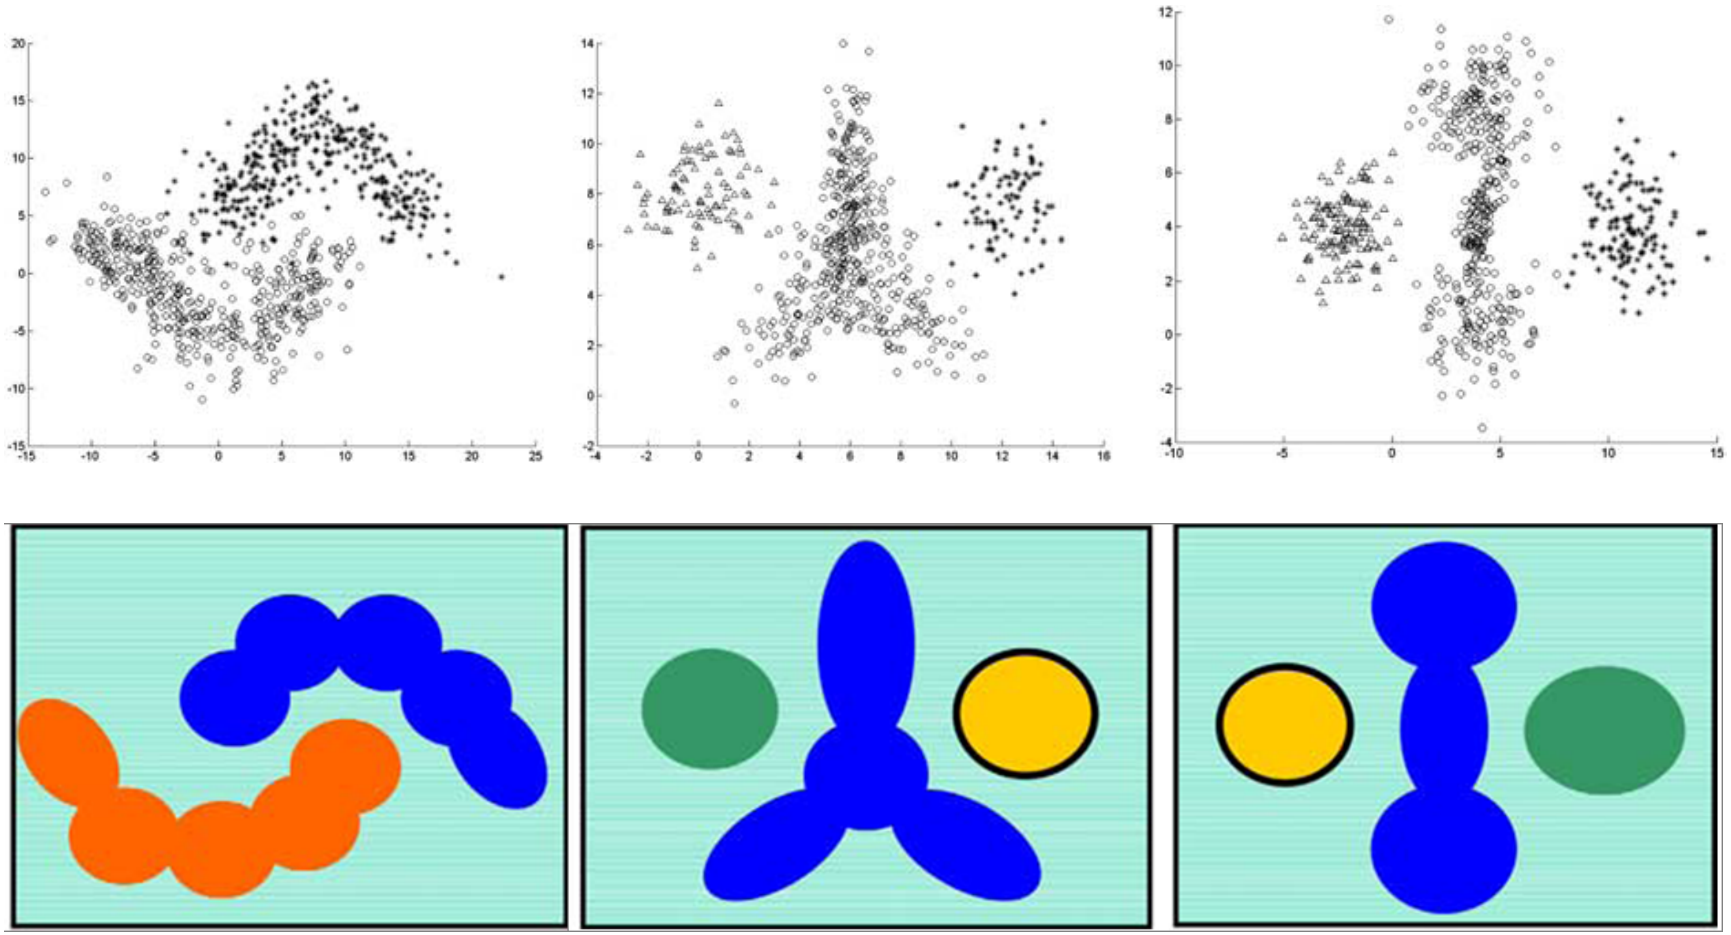

Examples of non-spherical clusters: (a) Teaeguk, (b) Triangle, (c) Xours (Cho et al.,2006)

Problems: local optima

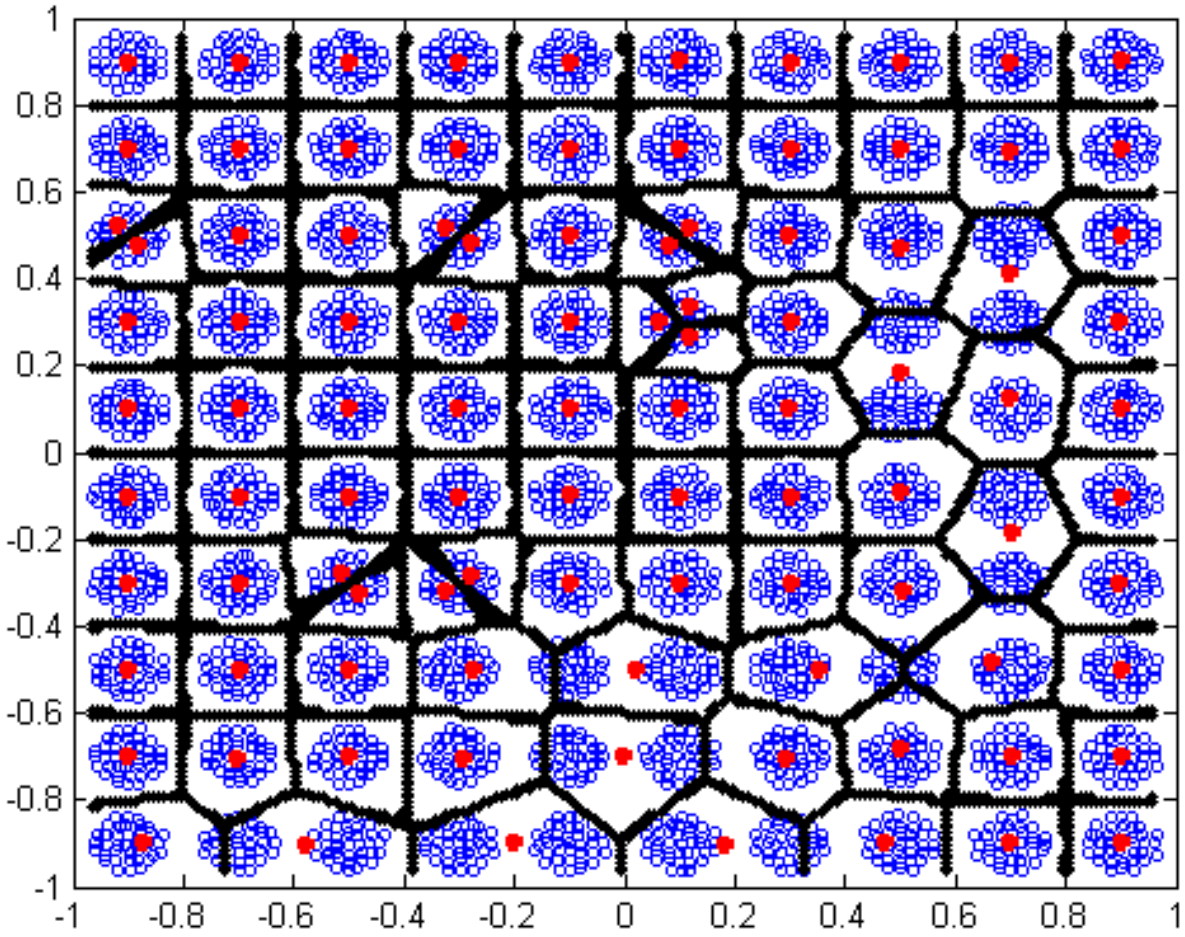

Checkerboard data with 100 data clusters and their cluster centers

Summary of concepts

- Clustering by distance to prototypes
- k-means clustering
- Quantization error
- Lloyd's algorithm
- Problems with Lloyd's algorithm
- Examples for the application of k-means clustering